

pc4_students.R

user

2020-03-19

```
#####  
## Course: Machine Learning for Economists and Business Analysts  
## Topic: High-Dimensional Confounding  
#####
```

```
#rm(list = ls())  
set.seed(100239)
```

```
#getwd()  
#setwd("")
```

```
# Load Packages  
library("fBasics")
```

```
## Warning: package 'fBasics' was built under R version 3.6.1  
## Loading required package: timeDate  
## Loading required package: timeSeries  
## Warning: package 'timeSeries' was built under R version 3.6.1
```

```
library("glmnet")
```

```
## Loading required package: Matrix  
## Loading required package: foreach  
## Loaded glmnet 2.0-18
```

```
library("AER")
```

```
## Warning: package 'AER' was built under R version 3.6.2  
## Loading required package: car  
## Warning: package 'car' was built under R version 3.6.2  
## Loading required package: carData  
## Warning: package 'carData' was built under R version 3.6.1
```

```
##
```

```
## Attaching package: 'car'
```

```
## The following object is masked from 'package:fBasics':
```

```
##
```

```
## densityPlot
```

```
## Loading required package: lmtest
```

```
## Loading required package: zoo
```

```
##
```

```
## Attaching package: 'zoo'
```

```

## The following object is masked from 'package:timeSeries':
##
##   time<-
## The following objects are masked from 'package:base':
##
##   as.Date, as.Date.numeric
## Loading required package: sandwich
## Loading required package: survival
library("hdm")

## Warning: package 'hdm' was built under R version 3.6.3
library("lmtest")
library("sandwich")
library("tidyverse")

## -- Attaching packages ----- tidyverse 1.3.0
## v ggplot2 3.2.0      v purrr  0.3.2
## v tibble  2.1.3      v dplyr 0.8.2
## v tidyr   0.8.3      v stringr 1.4.0
## v readr   1.3.1      v forcats 0.4.0

## -- Conflicts ----- tidyverse_conflict_
## x purrr::accumulate() masks foreach::accumulate()
## x tidyr::expand()      masks Matrix::expand()
## x dplyr::filter()      masks timeSeries::filter(), stats::filter()
## x dplyr::lag()         masks timeSeries::lag(), stats::lag()
## x dplyr::recode()      masks car::recode()
## x purrr::some()        masks car::some()
## x purrr::when()        masks foreach::when()

# Load data
df <- read.csv("job_corps.csv",header=TRUE, sep=",")

#####
### Exercise 1: Conventional Methods ###
#####

#####
# Task 1
#####

# Table with Descriptive Statistics
desc <- fBasics::basicStats(df) %>% t() %>% as.data.frame() %>%
  select(Mean, Stdev, Minimum, Maximum, nobs)
print(round(desc, digits=2))

##           Mean  Stdev Minimum Maximum  nobs
## EARNY4      204.44 195.69      0 2409.91 10516
## assignment    0.60  0.49      0    1.00 10516
## participation 0.44  0.50      0    1.00 10516
## female       0.43  0.49      0    1.00 10516
## age_1        0.41  0.49      0    1.00 10516
## age_2        0.31  0.46      0    1.00 10516

```

```
## age_3      0.27  0.45    0  1.00 10516
## ed0_6      0.26  0.44    0  1.00 10516
## ed6_12     0.36  0.48    0  1.00 10516
## hs_ged     0.24  0.43    0  1.00 10516
## white      0.26  0.44    0  1.00 10516
## black      0.49  0.50    0  1.00 10516
## hisp       0.17  0.38    0  1.00 10516
## oth_eth    0.07  0.26    0  1.00 10516
## haschld    0.20  0.40    0  1.00 10516
## livespou   0.06  0.24    0  1.00 10516
## everwork   0.80  0.40    0  1.00 10516
## yr_work    0.64  0.48    0  1.00 10516
## currjob    0.21  0.40    0  1.00 10516
## job0_3     0.22  0.41    0  1.00 10516
## job3_9     0.30  0.46    0  1.00 10516
## job9_12    0.21  0.41    0  1.00 10516
## earn1      0.11  0.31    0  1.00 10516
## earn2      0.27  0.44    0  1.00 10516
## earn3      0.14  0.34    0  1.00 10516
## earn4      0.07  0.25    0  1.00 10516
## badhlth    0.13  0.34    0  1.00 10516
## welf_kid   0.20  0.40    0  1.00 10516
## got_fs     0.45  0.50    0  1.00 10516
## publich    0.22  0.41    0  1.00 10516
## got_afdc   0.31  0.46    0  1.00 10516
## harduse    0.06  0.24    0  1.00 10516
## potuse     0.24  0.43    0  1.00 10516
## evarrst    0.24  0.43    0  1.00 10516
## pmsa       0.32  0.47    0  1.00 10516
## msa        0.47  0.50    0  1.00 10516
```

```
#####
# Task 2
#####
```

```
# Univariate OLS
```

```
ols <- lm(EARNY4 ~ participation, data = df)
summary(ols)
```

```
##
## Call:
## lm(formula = EARNY4 ~ participation, data = df)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -211.44 -168.41  -25.03  101.08 2210.97
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)   198.936     2.550   78.023 < 2e-16 ***
## participation  12.503     3.842    3.254  0.00114 **
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 195.6 on 10514 degrees of freedom
```

```
## Multiple R-squared:  0.001006,   Adjusted R-squared:  0.0009111
## F-statistic: 10.59 on 1 and 10514 DF,  p-value: 0.001141

# Robust standard errors
coeftest(ols, vcov = vcovHC(ols, type = "HC1"))

##
## t test of coefficients:
##
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)  198.9359    2.5047  79.4250 < 2.2e-16 ***
## participation  12.5026    3.8604   3.2387  0.001204 **
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

#####
# Task 3
#####

# IV results
iv <- ivreg(formula = EARNY4 ~ participation | assignment, data = df)
summary(iv)

##
## Call:
## ivreg(formula = EARNY4 ~ participation | assignment, data = df)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -213.84 -168.37  -23.14   100.57  2212.87
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)   197.042     3.058   64.434 < 2e-16 ***
## participation   16.803     5.428    3.096  0.00197 **
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 195.6 on 10514 degrees of freedom
## Multiple R-Squared:  0.0008871,   Adjusted R-squared:  0.0007921
## Wald test: 9.584 on 1 and 10514 DF,  p-value: 0.001968

# Robust standard errors
coeftest(iv, vcov = vcovHC(ols, type = "HC1"))

##
## t test of coefficients:
##
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)   197.0422     2.5047  78.6690 < 2.2e-16 ***
## participation  16.8027     3.8604   4.3526 1.358e-05 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

#####
### Exercise 2: Double Selection Procedure ###
#####
```

```
#####
# Task 1
#####

# Predict earnings
st1 <- rlasso(as.matrix(df[,c(4:ncol(df))]), as.matrix(df$EARNY4))
summary(st1)

##
## Call:
## rlasso.default(x = as.matrix(df[, c(4:ncol(df))]), y = as.matrix(df$EARNY4))
##
## Post-Lasso Estimation: TRUE
##
## Total number of variables: 33
## Number of selected variables: 16
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -384.45 -131.90  -19.95   90.80 2194.30
##
##              Estimate
## (Intercept)  189.994
## female      -56.801
## age_1        -3.582
## age_2         0.000
## age_3         8.514
## ed0_6         0.000
## ed6_12        0.000
## hs_ged       29.675
## white        22.872
## black       -24.628
## hisp          0.000
## oth_eth       0.000
## haschld       0.000
## livespou      0.000
## everwork     12.135
## yr_work      41.323
## currjob      18.280
## job0_3        0.000
## job3_9        0.000
## job9_12       5.634
## earn1       -25.530
## earn2         0.000
## earn3        24.463
## earn4        64.541
## badhlth       0.000
## welf_kid     -12.330
## got_fs       -6.239
## publich       0.000
## got_afdc      0.000
## harduse       0.000
## potuse      -16.297
```

```

## evarrst          0.000
## pmsa             0.000
## msa              0.000
##
## Residual standard error: 185.5
## Multiple R-squared:  0.1018
## Adjusted R-squared:  0.1005
## Joint significance test:
## the sup score statistic for joint significance test is 1818 with a p-value of 0

# Store selected variables
n1<- names(st1$coefficients[(st1$coefficients != 0) == TRUE])[-1]

#####
# Task 2
#####

# Predict participation
st2 <- rlasso(as.matrix(df[,c(4:ncol(df))]), as.matrix(df$participation))
summary(st2)

##
## Call:
## rlasso.default(x = as.matrix(df[, c(4:ncol(df))]), y = as.matrix(df$participation))
##
## Post-Lasso Estimation: TRUE
##
## Total number of variables: 33
## Number of selected variables: 1
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -0.4462 -0.4462 -0.4462  0.5538  0.6449
##
##              Estimate
## (Intercept)    0.446
## female         0.000
## age_1          0.000
## age_2          0.000
## age_3          0.000
## ed0_6          0.000
## ed6_12         0.000
## hs_ged         0.000
## white          0.000
## black          0.000
## hisp           0.000
## oth_eth        0.000
## haschld        0.000
## livespou      -0.091
## everwork       0.000
## yr_work        0.000
## currjob        0.000
## job0_3         0.000
## job3_9         0.000
## job9_12        0.000

```

```

## earn1          0.000
## earn2          0.000
## earn3          0.000
## earn4          0.000
## badhlth        0.000
## welf_kid        0.000
## got_fs          0.000
## publich         0.000
## got_afdc        0.000
## harduse         0.000
## potuse          0.000
## evarrst         0.000
## pmsa            0.000
## msa             0.000
##
## Residual standard error: 0.496
## Multiple R-squared:  0.002016
## Adjusted R-squared:  0.001921
## Joint significance test:
## the sup score statistic for joint significance test is 0.9382 with a p-value of 0.054

# Store selected variables
n2<- names(st2$coefficients[(st2$coefficients != 0) == TRUE])[-1]

#####
# Task 3
#####

# Take union of selected covariates
selected_covariates <- c("participation", unique(c(n1, n2)))

# Setup the formula of the linear regression model
sumx <- paste(selected_covariates, collapse = " + ")
linear <- paste("EARNY4",paste(sumx, sep=" + "), sep=" ~ ")
linear <- as.formula(linear)

# Post-Lasso OLS regression
ols <- lm(linear, data = df)
summary(ols)

##
## Call:
## lm(formula = linear, data = df)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -378.47 -132.12  -20.12   91.01 2201.29
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)    183.337     6.762   27.113 < 2e-16 ***
## participation    15.678     3.650    4.295 1.76e-05 ***
## female        -56.965     3.794  -15.015 < 2e-16 ***
## age_1          -4.133     4.565   -0.905 0.365348
## age_3           8.685     4.931    1.761 0.078229 .

```

```
## hs_ged      29.911      4.861      6.153 7.87e-10 ***
## white       23.023      5.167      4.456 8.44e-06 ***
## black      -24.811      4.509     -5.502 3.84e-08 ***
## everwork    12.168      6.236      1.951 0.051065 .
## yr_work     41.219      5.878      7.013 2.48e-12 ***
## currjob     18.076      5.020      3.601 0.000319 ***
## job9_12      5.685      5.675      1.002 0.316465
## earn1      -25.373      6.506     -3.900 9.68e-05 ***
## earn3       24.666      6.107      4.039 5.40e-05 ***
## earn4       65.047      8.651      7.519 5.99e-14 ***
## welf_kid    -12.366      4.794     -2.580 0.009903 **
## got_fs      -5.893      3.945     -1.494 0.135292
## potuse     -16.390      4.280     -3.830 0.000129 ***
## livespou    -2.070      7.570     -0.273 0.784519
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 185.5 on 10497 degrees of freedom
## Multiple R-squared:  0.1034, Adjusted R-squared:  0.1019
## F-statistic: 67.26 on 18 and 10497 DF,  p-value: < 2.2e-16
```

```
# Robust standard errors
```

```
coeftest(ols, vcov = vcovHC(ols, type = "HC1"))
```

```
##
## t test of coefficients:
##
##              Estimate Std. Error  t value  Pr(>|t|)
## (Intercept)  183.3367    6.8100  26.9219 < 2.2e-16 ***
## participation  15.6776    3.6616   4.2816 1.872e-05 ***
## female      -56.9647    3.7024 -15.3857 < 2.2e-16 ***
## age_1       -4.1325    4.6049  -0.8974 0.3695081
## age_3        8.6852    4.9891   1.7408 0.0817412 .
## hs_ged      29.9115    4.9989   5.9836 2.254e-09 ***
## white       23.0225    5.4843   4.1979 2.717e-05 ***
## black      -24.8106    4.5131  -5.4975 3.943e-08 ***
## everwork    12.1681    5.6126   2.1680 0.0301820 *
## yr_work     41.2189    5.5136   7.4758 8.286e-14 ***
## currjob     18.0755    5.3711   3.3653 0.0007673 ***
## job9_12      5.6854    6.4497   0.8815 0.3780699
## earn1      -25.3734    6.2281  -4.0740 4.655e-05 ***
## earn3       24.6663    6.8610   3.5952 0.0003257 ***
## earn4       65.0470   10.9807   5.9238 3.246e-09 ***
## welf_kid    -12.3665    4.6543  -2.6570 0.0078957 **
## got_fs      -5.8926    3.9438  -1.4941 0.1351761
## potuse     -16.3902    4.1337  -3.9650 7.389e-05 ***
## livespou    -2.0699    7.6720  -0.2698 0.7873174
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

```
#####
```

```
# Task 4
```

```
#####
```

```
# Double Selection Procedure
```



```

dsp <- rlassoEffect(as.matrix(df[,c(4:ncol(df))]), as.matrix(df$EARNY4)
, as.matrix(df$participation), model = TRUE, method = "double selection")
summary(dsp)

## [1] "Estimates and significance testing of the effect of target variables"
##      Estimate. Std. Error t value Pr(>|t|)
## d1      15.678      3.661    4.282 1.85e-05 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

#####
### Exercise 3: Double Machine Learning ###
#####

#####
# Task 1
#####
set.seed(1001)

# Partition Samples for Cross-Fitting
df_part <- modelr::resample_partition(df, c(obs_A = 0.5, obs_B = 0.5))
df_obs_A <- as.data.frame(df_part$obs_A) # Sample A
df_obs_B <- as.data.frame(df_part$obs_B) # Sample B

## Generate Variables
# Outcome
earnings_obs_A <- as.matrix(df_obs_A[,1])
earnings_obs_B <- as.matrix(df_obs_B[,1])

# Treatment
treat = 3 #Select treatment 2= offer to participate, 3 = actual participation
treat_obs_A <- as.matrix(df_obs_A[,treat])
treat_obs_B <- as.matrix(df_obs_B[,treat])


# Covariates
covariates_obs_A <- as.matrix(df_obs_A[,c(4:ncol(df_obs_A))])
covariates_obs_B <- as.matrix(df_obs_B[,c(4:ncol(df_obs_B))])

#####
# Task 2
#####

## Conditional Potential Earnings under Non-Participation
p = 1 # 1 for LASSO, 0 for Ridge
set.seed(100237)


## Using Sample A to Predict Sample B
# Potential Earnings under Non-Treatment
lasso_y0_A <- cv.glmnet(covariates_obs_A[treat_obs_A==0,], earnings_obs_A[treat_obs_A==0,],
                        alpha=p, type.measure = 'mse')
plot(lasso_y0_A)

```



```
fit_y0_A <- glmnet(covariates_obs_A[treat_obs_A==0,], earnings_obs_A[treat_obs_A==0,],
                  ,lambda = lasso_y0_A$lambda.min)
y0hat_B <- predict(fit_y0_A, covariates_obs_B)

## Using Sample B to Predict Sample A
# Potential Earnings under Non-Treatment
lasso_y0_B <- cv.glmnet(covariates_obs_B[treat_obs_B==0,], earnings_obs_B[treat_obs_B==0,],
                      alpha=p, type.measure = 'mse')
plot(lasso_y0_B)
```




```
fit_y0_B <- glmnet(covariates_obs_B[treat_obs_B==0,], earnings_obs_B[treat_obs_B==0,],
                  ,lambda = lasso_y0_B$lambda.min)
y0hat_A <- predict(fit_y0_B, covariates_obs_A, type = 'response')

#####
# Task 3
#####


## Propensity Score
p = 1 # 1 for LASSO, 0 for Ridge
set.seed(100236)

# Using Sample A to Predict Sample B
lasso_p_A <- cv.glmnet(covariates_obs_A, treat_obs_A, alpha=p, type.measure = 'mse')
plot(lasso_p_A)
```



```
fit_p_A <- glmnet(covariates_obs_A, treat_obs_A, lambda = lasso_p_A$lambda.min)
pscore_B <- predict(fit_p_A, covariates_obs_B)

# Using Sample B to Predict Sample A
lasso_p_B <- cv.glmnet(covariates_obs_B, treat_obs_B, alpha=p, type.measure = 'mse')
plot(lasso_p_B)
```



```
fit_p_B <- glmnet(covariates_obs_B, treat_obs_B, lambda = lasso_p_B$lambda.min)
pscore_A <- predict(fit_p_B, covariates_obs_A)

#####
# Task 4
#####

## Efficient Score
p_A = mean(pscore_A)
p_B = mean(pscore_B)

# Generate Modified Outcome in each sample
Y_star_A = invisible(treat_obs_A*(earnings_obs_A - y0hat_A)/p_A
  - (1-treat_obs_A)*pscore_A*(earnings_obs_A - y0hat_A)/(p_A*(1-pscore_A)))

Y_star_B = invisible(treat_obs_B*(earnings_obs_B - y0hat_B)/p_B
  - (1-treat_obs_B)*pscore_B*(earnings_obs_B - y0hat_B)/(p_B*(1-pscore_B)))

Y_star <- 0.5*(mean(Y_star_A) + mean(Y_star_B))

#####
# Task 5
#####

# Average Treatment Effect for Treated (ATET)
ATET <- round(Y_star, digits=3)
```

```

# Estimate variance for each sample
var_A = mean((Y_star_A - treat_obs_A*Y_star/p_A)^2)
var_B = mean((Y_star_B - treat_obs_B*Y_star/p_B)^2)

# Split sample estimator for standard error
SD_ATET <- round(sqrt(0.5*(var_A + (mean(Y_star_A) - Y_star)^2
                        + var_B + (mean(Y_star_B) - Y_star)^2)
                  /(length(Y_star_A) + length(Y_star_B))),digits=3)

print(paste0("Average Treatment Effect for Treated (ATET): ", ATET))

## [1] "Average Treatment Effect for Treated (ATET): 15.375"

print(paste0("Standard Error for ATET: ", SD_ATET))

## [1] "Standard Error for ATET: 3.663"

```